


# Resilience through chaos engineering

Azure Chaos Studio



# Cloud development



## The cloud has revolutionized application development

Management, maintenance, and security of physical infrastructure is eliminated.

PaaS services abstract away compute and network layers, enable faster development of complex applications.



## With this comes a lack of control

Outages have more impact and unknowns.

Disruptions to dependencies can have cascading effects.



## Applications must be designed for failure

Resilience is a shared responsibility between cloud provider and application developer.

Microsoft provides the Well Architected Framework and application development guidance.

**Resilience:** the capability of a system to handle and recover from disruption



Service outages impact availability



Availability impacts business: upsets customers + can lead to financial, life-and-death, or legal consequences



Outages steal developer time from working on improvements and new features




Quality practices need to be built-in to the entire service development and operation lifecycle

## Outages have major real-world consequences



2300  
flights

cancelled during Southwest  
Airlines datacenter outage,  
costing \$150M June 2016<sup>1</sup>




\$10  
refunds

after 30-minute YouTube TV  
outage during the 2018 World  
Cup, costing around \$8M<sup>2</sup>

2.5 hours

Walmart lost over \$9M in  
revenue during outage on  
Black Friday weekend 2018<sup>3</sup>



\$100M

Amazon lost almost \$100M in  
revenue in just 1 hour of  
downtime during Prime Day  
2018<sup>4</sup>




1. Travel and Leisure, [Southwest computer outage cancels flights](#), 2016 2. Business Insider, [J. Crew website crashes on Black Friday](#), 2018 3. Variety, [YouTube TV Offers one week credit after World Cup outage](#), 2018 4. Business Insider, [Amazon Prime Day Website Issues](#), 2018

# What is chaos engineering?

Improve resilience with systematic controlled chaos

- In the cloud, dependencies can be disrupted at any time
- Design for failure, architect for resilience
- Validate resilience through fault injection to confirm that your solution can withstand the disruptions it will encounter in production



## Chaos **use cases**

Perform automated and UI-driven resilience validation

- Disrupt dependencies with ad hoc experiments in a dev/test environment to validate new code
- Gate code flow in CI/CD pipeline automation
- Perform incident fix and incident regression testing
- Host a drill event or game day
- Validate on call and livesite process
- Conduct Error Budget testing

Leverage the **scientific method** to formulate **hypotheses** around resiliency **scenarios**, choose a safe environment for experimentation, craft and execute a fault injection **experiment**, **monitor** the impact, analyze results and make improvements.

## Azure Chaos Studio

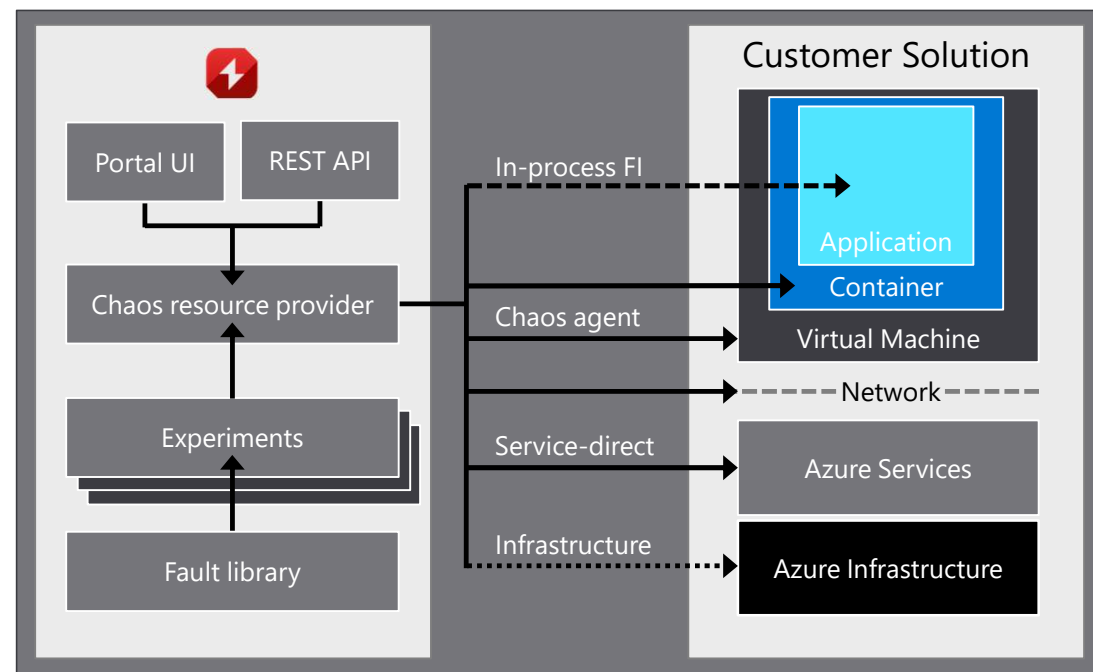
Measure, understand, improve, and maintain product resilience

- Fully [managed service](#) to validate Azure application and service resilience.
- Deep [Azure integration](#), including Azure Portal, Azure Resource Manager, and Azure Monitor.
- [Portal UI](#) and [REST API](#) and [SDKs](#) to execute manual and automated chaos experiments.
- [Expandable library](#) of common resource pressure and dependency disruption faults and actions.
- Advanced [workflow orchestration](#) for manual and automated fault injection experiments that simulate [real-world scenarios](#) with parallel and sequential action execution.
- [Safeguards](#) to minimize impact radius and to control experiment execution.

# Azure Chaos Studio

## Measure, understand, improve, and maintain product resilience

- Fully [managed service](#) to validate Azure application and service resilience
- Deep [Azure integration](#), including Azure Portal, Azure Resource Manager, Azure Monitor
- [Portal UI](#) and [REST API](#) and [SDKs](#) to execute manual and automated chaos experiments
- [Expandable library](#) of common resource pressure and dependency disruption faults and actions
- Advanced [workflow orchestration](#) for manual and automated fault injection experiments that simulate [real-world scenarios](#) with parallel and sequential action execution
- [Safeguards](#) to minimize impact radius and to control experiment execution





# Disruption types

## Expanding library of faults and actions across the entire Azure stack



### Chaos for Applications

In-process fault injection to perturb managed code function and API calls


- Azure SDKs (*preview*)
- Customer code (*future*)



### Guest OS agent

Chaos agent for Windows and Linux VMs


- Resource pressure faults (CPU, memory, disk, network)
- Dependency disruption faults (process/service, network)



### Service direct

Agentless and mock/proxy faults against Azure services

- Dependency disruptions
- Configuration changes
- Latency



### Chaos for Infrastructure

Used internally by Microsoft teams to disrupt Azure infrastructure

- Chaos agent deploy to host
- Used by Microsoft to improve the Azure platform

# Chaos **experiment**

Orchestrated, real-world scenarios with fault actions applied to resource targets while under load

## **Elements**

### **Application**

Application/service for resilience validation.

### **Target**

Resource target(s) enabled for use in chaos experiments.

### **Workload**

Synthetic workload or real customer traffic for representative customer usage.

### **Fault actions**

Orchestrated fault and other (delay, load) actions.

### **Observability**

Telemetry and thresholds to monitor application health.

## **Process**

### **Formulate hypothesis**

What scenario is being validated and what are possible and expected outcomes?

### **Craft experiment**

Orchestrate workload execution with fault actions against subscription resource targets.

### **Execute and monitor**

Compare baseline observability metrics, monitoring telemetry, reporting and analysis.

### **Analyze and improve**

New code, code and configuration changes.

### **Steady state**

Rinse and repeat; add continuous production monitoring and validation.

# Chaos Studio use cases

## Scenario Validation

- Availability Zone down
- DNS outage
- AAD outage
- Region down
- Region failover
- BC/DR and HA/DR
- Spike/crush load (Black Friday)
- Network outage
- Resource pressure and noisy neighbor
- Emergency certificate rotation
- Storage failover
- Autoscale validation
- Systematic dependency disruption
- Maintenance events
- and much more...

## Fault and Action Library

Agent (Windows, Linux)	Service-direct
<ul style="list-style-type: none"><li>• CPU pressure</li><li>• Physical memory pressure</li><li>• Virtual memory pressure</li><li>• Disk I/O pressure</li><li>• Kill process</li><li>• DNS failure (Windows)</li><li>• Network latency (Windows)</li><li>• Network disconnect (Windows)</li><li>• Network firewall disconnect (Windows)</li><li>• Stress-ng stress (Linux)</li><li>• Stop service (Windows)</li><li>• Change time (Windows)</li></ul>	<ul style="list-style-type: none"><li>• VM shutdown or kill</li><li>• VMSS shutdown or kill</li><li>• Classic Cloud Services shutdown</li><li>• Cosmos DB failover</li><li>• Azure Cache for Redis reboot</li><li>• Network security group set rule</li><li>• Key Vault deny access</li><li>• Key Vault increment certificate version</li><li>• Key Vault disable certificate</li><li>• Key Vault update certificate policy</li><li>• Disable autoscale</li><li>• AKS network disconnect/packet loss/latency</li><li>• AKS pod failure, container failure/kill</li><li>• AKS CPU, memory pressure</li><li>• AKS file I/O latency/failure</li><li>• AKS DNS failure</li><li>• AKS time change</li><li>• AKS HTTP delay, modification</li></ul>

# Chaos Studio Roadmap

[Cu]

- ✓ Support for chaos experiments on customer Private Networks
- ✓ Dynamic Target selection for VMSS Availability Zones
- ✓ In-process Fault Injection pilot
- ✓ Deployment to more regions
- ✓ Faults
  - ✓ Azure Cache for Redis reboot
  - ✓ Classic Cloud Services shutdown
  - ✓ Key Vault Deny Access

[Zn]

- Dynamic Target selection across more resource properties
- User assigned managed identity
- Identity Management improvements
- Chaos agent Private Network Support
- ADO Pipeline task for automation
- Observability integration and hypothesis attainment scoring
- In-Process Fault Injection preview
- Azure Load Test Integration
- Infrastructure multi-step, multi-branch
- Faults
  - Disable, Update, Increment Cert
  - Disable Autoscale
  - Network packet drop (Windows agent)
  - Host suspend/resume

[Ga]

- Drill Manager UX
- In-Process Fault Injection UX
- Observability stop conditions
- Resiliency Score
- Canary Drills powered by Chaos Studio
- More faults...
- More regions...
- Billing
- GA fall 2023

Get started at <https://aka.ms/AzureChaosStudio>

# Get started

<http://aka.ms/AzureChaosStudio>

1. **Explore** – Become familiar with the Chaos Studio user interface. Set up a simple VM, enable agent and service-direct faults against it, create and execute an experiment to shutdown the VM or apply CPU pressure, and monitor the impact.
2. **Pilot** – Integrate chaos provisioning and agent deploy into a test environment and add chaos experiments to integration, stress, or other tests. Gauge impact through existing monitoring and impact to established baselines.
3. **Automate** – Provision chaos and agent deploy as part of test environment buildout or daily automation pipeline. Add automated chaos experiment execution to ADO Pipeline or CI/CD process in a pre-production environment to catch issues before they impact customers.
4. **Drill** – Perform a drill or game day event. Choose a scenario, establish a hypothesis or hypotheses, set aside a day, author and pre-test an experiment or experiments, choose participants and preview monitoring, perform the drill on the chosen day, monitor closely, analyze results, create repair items, rinse and repeat.

## FAQ

1. Q: Billing?  
A: Eventually. Planning pay-as-you-go model.
2. Q: Do you support faults against on-premise or cross-cloud resources? how about Azure Edge or Arc?  
A: No. Network faults can be used to block access to these, but at this time we do not support non-Azure and on-premise resources. This is in our roadmap.
3. Q: Does Chaos Studio support rollback at the end of an experiment or if things go wrong?  
A: No. At end of experiment, non-destructive fault behaviors end (CPU, memory, disk, network pressure) and systems return to pre-experiment state but if a VM or process is killed, Cosmos DB failed over, etc. nothing is done to restore state.
4. Q: Are you actually *causing* or *simulating* issues and outages?  
A: Both... it depends on both the fault and the environment. Faults can be combined to represent more complex scenarios.
5. Q: Can I create my own fault?  
A: Coming soonish: BYOF (internal/external) and inner source (internal)
6. Q: Does Chaos Studio create a workload or load against my application?  
A: No. Chaos Studio is primarily a platform for orchestrating fault injection. See ALT. Coming soon: ALT load action. Coming later: custom load action
7. Q: Do you have PaaS fault coverage for services XXX, YYY, and ZZZ?  
A: Not yet, but the NSG and network faults can impact many services

# Demo



Azure Chaos Studio

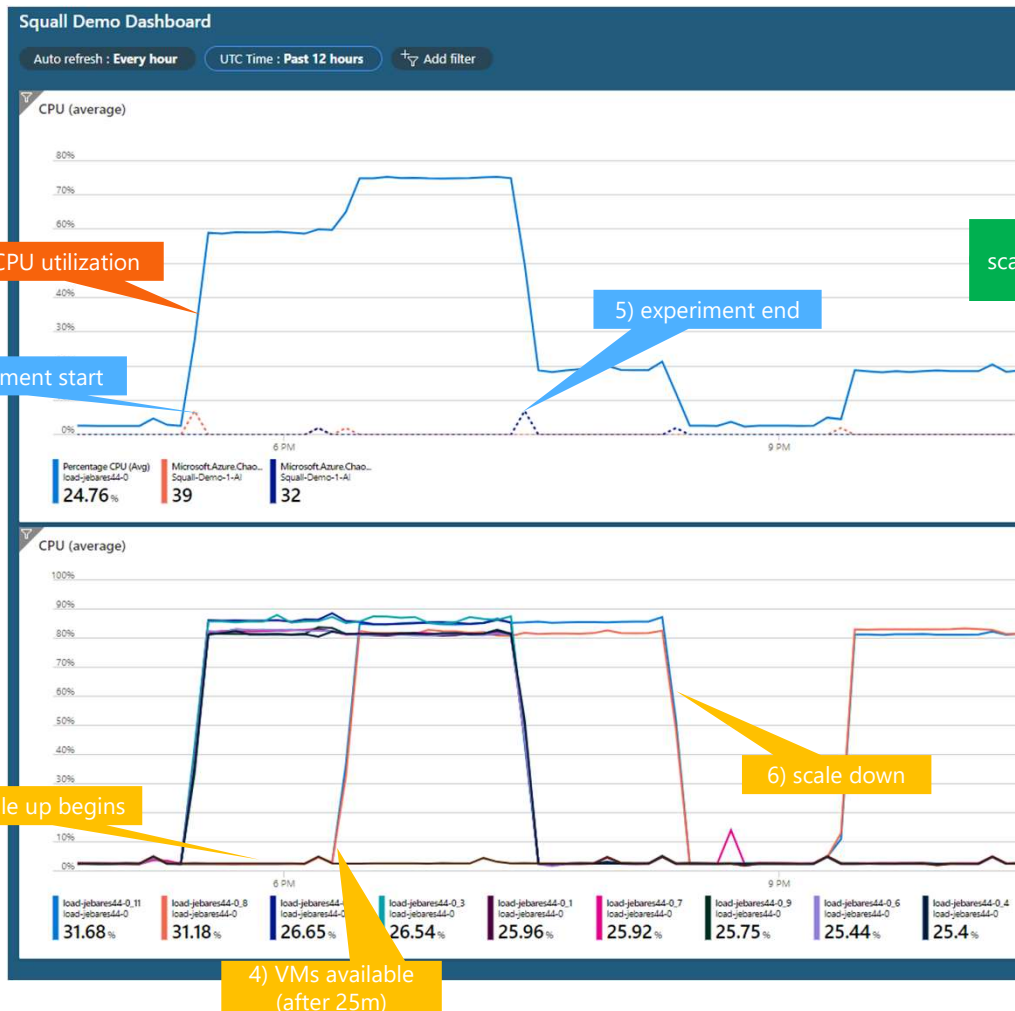


Thank you.



# Case Study: autoscale validation

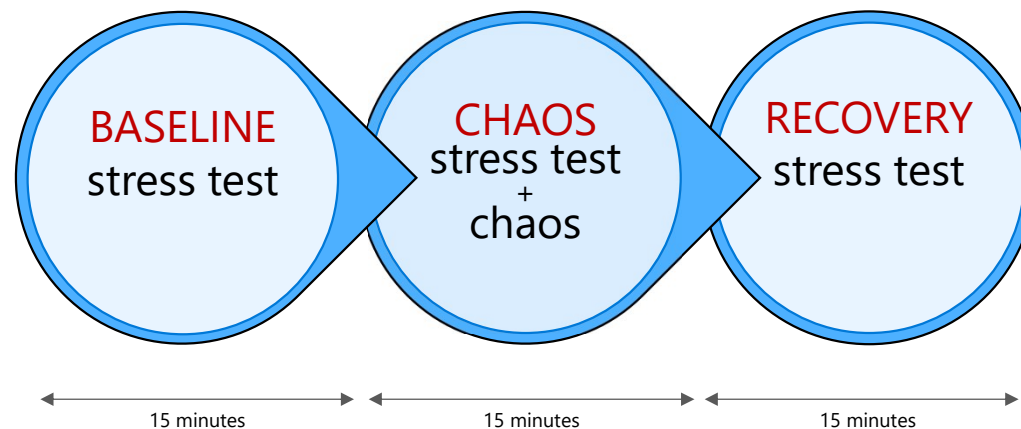
- Issue: machines provisioned, but AGS not deployed (autoscale worked, got billed, but service not deployed)



# Case Study: nightly validation of new builds

Daily build deployed to VM test cluster.

Chaos added to stress test pass, stress extended to 3 phases with resource pressure faults (CPU, Memory, Disk IO, etc.) run across 10 VMs.



Multiple metric comparisons enabled

- **Baseline** - compare throughput and other test metrics against goals and historical trends
- **Chaos** - measure the impact of chaos to SLO and service health
- **Recovery** - validate time to return to steady state